

Validating Leon's Inevitable Business Formula

Executive summary

Leon's formula can be treated as a screening rule for redress businesses: a market becomes most attractive when a bad event happens to users without their fault, another party is objectively required to compensate or fix it, and the user can trigger that remedy with very little effort. In Leon's own framing, the pattern is "not your mistake + someone owes you + easy action." ¹

I did **not** find a peer-reviewed literature stream that uses the exact label "inevitable business." The strongest support is **indirect but strong** and comes from three bodies of evidence that fit the formula closely: service failure and service recovery research, administrative burden and sludge research, and official legal or contractual remedy frameworks such as EU passenger rights, U.S. warranty law, regulated utility compensation, and service-level agreements. ¹

That means the formula is best understood not as a universal business law, but as a **high-signal heuristic for claims and enforcement markets**. It is strongest where entitlement is objective, evidence is easy to verify, user blame is low, and the remedy can be turned into a short flow. This is exactly why businesses like AirHelp and Flightright exist and why regulated schemes like Ofcom's automatic compensation work so well. ²

A practical validation system should score four blocks: **failure inevitability**, **obligation strength**, **user simplicity**, and **economics**. A useful greenlight rule is: total score at least 70 out of 100, no critical block below 60 percent of its own maximum, and confidence at least 65 out of 100. If an idea fails the obligation block or the simplicity block, it is usually not an inevitable business, no matter how painful the problem is.

Concept and testable hypotheses

Leon's formula becomes usable when each component is turned into a testable hypothesis.

Unavoidable failure should mean two things. First, the user is not the main cause of the event. Second, the event happens often enough, and persists enough, to create repeatable demand. The service-failure literature supports this direction: customers react strongly to failures, and complaint handling quality affects trust and commitment. A recent airport study also found that blame attribution matters, with internally caused delays hurting satisfaction more than weather-related delays that customers tend to excuse. ³

Obligation to pay or fix should mean there is a formal duty, not just a moral preference. The strongest versions are statutes, regulations, tariffs, warranties, and contractual SLAs with clear triggers, clear exceptions, and defined remedies. EU passenger-rights law is a textbook example because the carrier owes compensation, care, and re-routing in defined cases, while also carrying the burden to justify extraordinary

circumstances. U.S. warranty law is another strong example because it requires timely remedy without charge and, after failed repair attempts, a refund or replacement option. ⁴

Simple user action should mean that a customer can start the process quickly with few inputs, few uploads, and a high completion rate. This is where administrative burden and sludge matter most. That literature argues that valid entitlements often go unused because of learning costs, psychological costs, and compliance costs. Businesses like AirHelp and Flightright exist precisely because they reduce those process costs for already-valid claims. ⁵

These ideas can be converted into a direct hypothesis table:

Component	Research question	Pass hypothesis
Failure inevitability	Does the triggering event happen often enough, with low user culpability, and stay persistent in the market?	Event incidence is meaningful, user-caused share is low, and residual event rate stays above your minimum viable threshold
Obligation strength	Is there a real duty to compensate, repair, or reimburse?	Law, regulation, tariff, warranty, or SLA defines trigger, remedy, and exceptions
Simple user action	Can the user complete the first claim step with very low effort?	Median initial flow under 3 minutes, low field count, minimal documents, high mobile completion
Economics	Can acquisition, verification, case handling, and collection produce positive unit economics?	Verified eligible cases generate positive contribution margin and acceptable cash cycle

The most important implication is simple: **pain alone is not enough**. A market only starts to look inevitable when pain is combined with enforceability and low-friction action. ⁶

What the literature and the law support

The strongest academic support comes from **service failure and recovery**. Tax, Brown, and Chandrashekar found that customers judge complaint experiences through outcome, procedure, and interpersonal treatment, and that poor complaint handling weakens trust and commitment. Smith, Bolton, and Wagner found that customers prefer recovery resources that match both the type and magnitude of the failure. Those findings fit Leon's logic well because they show why a third party can create real value by organizing recovery after a failure has already happened. ⁷

The next support comes from **administrative burden and sludge**. The literature summarized in this stream treats friction as a real suppressor of entitlement take-up. Burden et al., Heinrich, Herd and Moynihan, and sludge-oriented work associated with Thaler and Sunstein all point in the same useful direction for product teams: extra forms, unclear steps, repeated document requests, and waiting costs are not minor UX issues, they are core determinants of whether people actually claim what they are owed. ⁸

The legal support is especially strong. EU Regulation 261/2004 creates objective rights for denied boarding, cancellation, delay, care, notice, and reimbursement. The regulation states that compensation applies in

defined cases, and the carrier must prove extraordinary circumstances to avoid payment. It also requires airlines to inform passengers of their rights. U.S. Magnuson-Moss creates federal minimum standards for written warranties, including timely remedy without charge and a refund or replacement after a reasonable number of failed repair attempts, while limiting extra duties imposed on consumers. These are exactly the kinds of rule systems that can be translated into a business. ⁹

Recent market evidence reinforces the same point. AirHelp says it has helped 28 million passengers, paid over 3 million passengers, offers a 2-minute claim check, and charges a standard 35 percent fee including VAT. Flightright reports over €700 million in compensation paid out, more than 12 million passengers supported, a 2-minute check, 99 percent success rate claims, and a no-win-no-fee structure. These companies are not proof of the formula in the academic sense, but they are strong market evidence that legally valid but operationally difficult claims can support scalable businesses. ¹⁰

A final important boundary comes from **automatic compensation**. Ofcom's broadband and landline scheme pays users without them needing to ask, provided their provider is part of the scheme and the trigger conditions are met. This shows that the formula works operationally even more strongly when friction becomes almost zero. But it also shows a limit for startups: when the regulator or provider already automates the remedy, the room for a new intermediary shrinks sharply. ¹¹

Data model and measurement system

A product team should validate the formula with four data layers: **event frequency**, **obligation structure**, **user-effort funnel**, and **economics**.

For **event frequency**, start with public event and complaint data. In aviation, BTS publishes airline on-time performance and related traffic data, including Airline Service Quality Performance 234. In consumer-financial complaints, CFPB publishes a downloadable database and an API, while also warning that complaint volume must be interpreted with company size and market share in mind. In insurance, NAIC provides consumer complaint and company-search tools. For legal sources, EUR-Lex and ELI resources are especially useful because they support identifiers, ontology-based metadata, XML encoding, authority tables, and validation tools. ¹²

For **obligation structure**, create one machine-readable rule record per jurisdiction and product type. A practical schema should include: jurisdiction, source type, trigger event, trigger threshold, obligor, remedy type, remedy formula, evidence required, exceptions, burden of proof, and time limit. This matches how official legal sources express real obligations. EU passenger-rights sources define compensation brackets and extraordinary-circumstance logic. Magnuson-Moss defines repair, refund, replacement, and consumer-duty limits. AWS EC2 SLA defines availability thresholds, credit percentages, required logs, claim windows, and exclusions. ¹³

For **simple user action**, measure the claim flow as a funnel, not by intuition. Use these metrics:

Metric	How to measure	Suggested internal pass threshold
Time to first value	seconds until user sees likely eligibility result	under 60 seconds
Mandatory user-entered fields	count of required fields in initial screen	8 or fewer
Document burden before eligibility result	number of required uploads	0 to 1
Claim-start completion	completed starts ÷ started flows	at least 60%
Median initial-flow time	minutes for eligible first-time user	under 3 minutes
Ease score	post-flow survey, 1 to 7 scale	at least 5.5
Human-support dependency	share requiring live support at screening stage	below 30%

These are **decision thresholds for validation**, not published market norms. They are meant to translate burden/sludge theory into a product operating standard. ¹⁴

For **unavoidable versus reducible failures**, use a two-part model. First, classify whether the user caused the event: user-caused, mixed, provider-caused, third-party-caused, or unknown. Leon's framing heavily favors "not the user's fault." ¹⁴

Second, classify reducibility at the system level:

- **Persistent residual failure:** still common after controlling for provider quality, route or product mix, seasonality, and obvious fraud.
- **Shrinking structural failure:** falls rapidly as providers improve operations.
- **Already automated resolution:** provider or regulator already compensates automatically.

Ofcom's scheme is the clearest example of the last category. It still proves the formula's logic, but usually weakens the business opportunity for a standalone intermediary. ¹¹

For **economics**, track: eligible incident volume, claimable value per incident, lead-to-verified eligibility rate, approval rate, time to payout, CAC, support minutes per case, legal-escalation rate, denial rate, and contribution margin per paid case. Public company-style signals from AirHelp and Flightright are useful directional benchmarks for economics and process design, but your own economics test should be the decision driver. ¹⁵

Validation protocol and scoring rubric

Use a stage-gated validation process. The goal is to kill weak ideas fast.

Discovery gate

In the first pass, collect one official rule source, one operational guidance source, one event-data source if available, and one market-activity signal such as a complaint registry, payout scheme, or claims competitor. If you cannot identify a real trigger, a real obligor, and a real remedy formula quickly, fail the idea. Markets with unclear duties almost always break later. ¹⁶

Measurement gate

Score the idea with this rubric:

Block	Weight	What to score
Failure inevitability	30	frequency, user non-culpability, persistence, objective detectability
Obligation strength	35	legal or contractual force, remedy clarity, exception clarity, burden-of-proof position
User simplicity	25	time, field count, document burden, completion, need for support
Economics	10	paid-case margin, CAC realism, payout cycle, operational complexity

Use the following pass rule:

- Total score at least **70 / 100**
- Failure inevitability at least **18 / 30**
- Obligation strength at least **21 / 35**
- User simplicity at least **15 / 25**

If the market fails obligation strength or user simplicity, reject it even if the total score is close.

Confidence score

Add a separate confidence score from 0 to 100:

- 40 points for data quality and freshness
- 30 points for causal evidence quality
- 20 points for legal clarity
- 10 points for benchmark comparability

Interpret it like this:

- **80 to 100**: high confidence
- **65 to 79**: medium confidence
- **below 65**: not enough evidence for launch

Sample sizes and tests

For rapid experiments, use standard two-arm designs.

For binary outcomes such as claim-start completion:

- Detecting **20% to 30%** improvement requires about **294 users per arm**
- Detecting **35% to 45%** improvement requires about **376 users per arm**

For continuous effort metrics:

- Detecting **0.4 standard deviations** requires about **99 users per arm**
- Detecting **0.3 standard deviations** requires about **175 users per arm**

Use:

- two-sample proportion tests for completion and conversion
- logistic regression for eligibility and approval prediction
- Welch t-tests or linear models for effort scores and time-to-complete
- survival analysis for time-to-payout and time-to-resolution
- difference-in-differences or natural experiments when regulation changes or providers adopt automation

A practical sequence is: 1. Measure event frequency and legal rule clarity. 2. Build a fake-door or MVP screening flow. 3. Run friction experiments. 4. Run assisted manual claims for a small cohort. 5. Estimate eligibility, approval, payout time, and margin. 6. Re-score the market.

Suggested pass-fail output

Your system should produce both a decision and an explanation:

```
{
  "industry": "air passenger compensation",
  "jurisdiction": "EU",
  "decision": "PASS",
  "total_score": 84,
  "confidence_score": 78,
  "scores": {
    "failure_inevitability": 24,
    "obligation_strength": 32,
    "user_simplicity": 20,
    "economics": 8
  },
  "key_reasons": [
    "Objective statutory trigger and payout brackets",
    "Carrier bears burden on extraordinary circumstances",
    "User action can be reduced to a short screening flow"
  ]
}
```

```

],
"red_flags": [
  "Evidence requirements vary by airline and country",
  "Possible future legal reform risk"
],
"recommended_next_test": "Run friction A/B test on 2-minute claim-start flow"
}

```

Prototype experiments and benchmark cases

The fastest way to validate a candidate market is to run **small, high-signal experiments** before building full automation.

A useful MVP stack is:

- a landing page that states the trigger and promised outcome
- a 2-minute claim eligibility flow
- live human verification behind the scenes
- a low-risk fee offer such as no-win-no-fee
- a simple document upload step only after likely eligibility

Use these experiment types:

Experiment	What it tests	Primary metric	Good early signal
Eligibility landing page	latent demand	visitor to flow-start rate	at least 8%
Fee framing test	willingness to delegate	flow start to case submit	no-win-no-fee outperforms fixed-fee
Friction manipulation	sensitivity to effort	completion rate	every removed field lifts completion
Evidence timing test	when to ask for documents	submit rate and fraud rate	later request improves submit without large fraud cost
Manual concierge test	true approval and payout	paid-case rate	enough to estimate unit economics
Escalation path test	legal process tolerance	share requiring court or regulator escalation	low enough to keep ops margin healthy

Those thresholds are **internal validation rules**, not sourced industry averages.

A useful friction experiment is to randomize one extra burden at a time:

- ask for ID before eligibility vs after eligibility

- require flight or contract number up front vs optional
- one-page flow vs multi-step wizard
- show expected payout before entering details vs after
- mobile camera upload vs file picker only


If the market is truly “inevitable,” completion should remain relatively strong under realistic evidence requirements, but should still improve materially when burdens are removed. That pattern is exactly what burden/sludge logic predicts. ¹⁴

Benchmark case comparison

Case	Failure type	Obligation source	User action complexity	Why it fits or fails
AirHelp	flight delay, cancellation, denied boarding	EU passenger-rights law and related regimes	low, 2-minute check, handled end-to-end	strong fit, because legal duty exists and the company removes hassle; AirHelp reports 28M+ passengers helped and 3M+ paid ¹⁷
Flightright	flight delay, cancellation, replacement travel reimbursement	EU261 and related rights	low to medium, claim flow simplified, no-win-no-fee	strong fit, because remedy is formulaic and user delegation is easy; Flightright reports €700M+ paid and 12M+ passengers supported ¹⁸
Warranty services	product defect or malfunction	Magnuson-Moss plus state warranty law	medium, can be simple if product and purchase are identifiable	moderate fit; obligation is real, but failure frequency and proof quality vary by category ¹⁹
Utility or telecom refunds	loss of service, missed appointment, delayed activation	regulator code or tariff	very low when automatic	strong formula logic, but weaker startup opportunity if compensation is already automatic, as in Ofcom's scheme ¹¹
Cloud SLA credits	uptime failure	contract SLA	low to medium, often log-based and formulaic	good B2B fit where observability is strong; weaker if manual claim process or exclusions dominate ²⁰

Case	Failure type	Obligation source	User action complexity	Why it fits or fails
Insurance claims assistance	covered loss or denied claim	policy contract plus regulated claims processes	medium to high	mixed fit; obligation exists, but proof burden, adjudication, and discretion are often much heavier. NAIC consumer tools confirm complaints and search tools exist, but the workflow is less inherently simple ²¹

Suggested visualization



Custom GPT implementation

OpenAI's official GPT builder supports instructions, knowledge files, capabilities, testing in Preview, and actions; knowledge is meant for source material, while instructions define behavior. OpenAI also recommends explicit structure, examples, clear headings, and test/evaluation loops. GPTs can attach up to 20 files, with each file up to 512 MB. ²²

GPT role

The GPT should act as a **business validator for entitlement and recovery markets**. It should not hype ideas. It should accept an industry as input and produce a scored decision with evidence, thresholds, and open questions.

Required system behavior

Use instructions that force the GPT to:

- identify the failure event
- identify the obligor
- locate official rule sources first
- classify user culpability
- score friction in the claim flow
- estimate which parts are objective vs discretionary
- output structured tables and a pass-fail decision
- separate evidence from inference
- highlight what is missing

Recommended prompt structure

Following OpenAI's guidance, use a structured instruction with identity, instructions, examples, and context.

23

A strong Custom GPT instruction skeleton:

```
# Identity
You are a research and product-validation assistant for "inevitable business"
ideas.
Your job is to test whether a market fits this formula:
Inevitable business = unavoidable failure + obligation to pay/fix + simple user
action.

# Instructions
- Always begin by identifying:
  1. Failure event
  2. User culpability
  3. Obligated party
  4. Source of obligation
  5. Trigger threshold
  6. Remedy type
  7. Required user action
- Prioritize official legal, regulatory, and company sources.
- Score the market on:
  - Failure inevitability
  - Obligation strength
```

- User simplicity
- Economics
- Produce:
 - executive summary
 - evidence table
 - scoring table
 - pass/fail decision
 - confidence score
 - open questions
- Clearly mark any inference.
- Never present unsupported claims as facts.

Output format

Return:

1. Summary
2. Scoring table
3. Rule extraction table
4. Recommended tests
5. Final decision JSON

Examples

<example_input>

Industry: airline delay compensation in the EU

</example_input>

<example_output>

Decision: PASS

Reason: strong legal entitlement, objective trigger, high claimant friction that can be reduced

</example_output>

Input schema

```
{
  "industry": "string",
  "jurisdiction": "string",
  "customer_type": "consumer|SMB|enterprise",
  "failure_event": "string",
  "current_user_flow": "string",
  "known_obligor": "string|null",
  "known_rule_sources": ["string"],
  "assumptions": ["string"]
}
```

Output schema

```
{
  "summary": "string",
  "score_total": 0,
  "confidence_score": 0,
  "subscores": {
    "failure_inevitability": 0,
    "obligation_strength": 0,
    "user_simplicity": 0,
    "economics": 0
  },
  "rule_extraction": [
    {
      "jurisdiction": "string",
      "source_type": "string",
      "trigger_event": "string",
      "trigger_threshold": "string",
      "obligor": "string",
      "remedy_type": "string",
      "remedy_formula": "string",
      "exceptions": ["string"],
      "burden_of_proof": "string"
    }
  ],
  "decision": "PASS|FAIL|UNCERTAIN",
  "open_questions": ["string"],
  "recommended_next_tests": ["string"]
}
```

Example user prompts for the GPT

- "Evaluate whether missed package-delivery compensation is an inevitable business in Germany."
- "Score EV charging downtime refunds in California against the formula."
- "Compare broadband outage compensation in the UK vs the US."
- "Give me a pass/fail on appliance warranty claim automation in Texas."

Evaluation prompts for the GPT

- "Which claim in your report depends on inference rather than a primary source?"
- "What evidence would most likely change your decision?"
- "Show me the friction points that most reduce take-up."
- "Convert your rule extraction into a machine-readable table I can load into a database."

Suggested presentation style

Have the GPT always present:

- one short executive summary
- one scoring table
- one rule extraction table
- one funnel table for user effort
- one mermaid decision diagram when useful
- one final JSON block for downstream systems

Risks and open questions

The biggest limitation is conceptual. Leon's formula is **strongest in claims and entitlement markets**, but it is not a complete model of all durable businesses. A product can be huge without fitting this structure, and a market can fit the structure yet still be unattractive because it is shrinking, already automated, or too compliance-heavy. ²⁴

There is also a legal-risk issue. These businesses depend on continuing access to enforceable remedy frameworks. If a regulator automates compensation, your role may disappear. If a law is weakened, thresholded upward, or narrowed by exceptions, case volume and economics can break. This is especially important in areas like passenger rights, where reform has been debated and operational interpretation can vary. ²⁵

A second limitation is data quality. Complaint and public-registry data are useful, but they are not perfect demand measures. CFPB explicitly warns that complaint data is not a statistical sample and that complaint counts should be interpreted with market-share and population context. The same caution should apply across industries. ²⁶

A third limitation is ethics and compliance. Reducing friction is good for users, but firms in this category often process ID documents, payment data, travel data, damage reports, and dispute records. That means fraud controls, privacy controls, explainable decisions, and clear fee disclosure are not optional. Warranty and passenger-rights laws also include disclosure duties and consumer-information obligations that should be reflected in product design. ²⁷

The most important open research questions are:

- Which industries have high legal entitlement but still poor take-up because of friction?
- Which categories are being made automatic by regulators, shrinking startup space?
- Where can proof be verified from third-party logs rather than from user memory?
- Which markets have enough claim value per case to offset support and compliance costs?

Recommended eight-week research sprint

A practical sprint can be run in eight weeks.

Week	Goal	Deliverable
1	Build candidate list by industry and jurisdiction	longlist of 20 markets
2	Extract official obligations	rule database with trigger, remedy, exception, burden fields
3	Estimate event frequency and complaint volume	demand memo with sources and caveats
4	Design and ship fake-door or MVP screening flow	working landing page and analytics dashboard
5	Run first friction A/B tests	completion and effort report
6	Run manual concierge claims on small cohort	verified eligibility and payout report
7	Model unit economics and re-score top markets	pass-fail ranking with confidence scores
8	Build Custom GPT and evaluator prompts	reusable GPT plus decision JSON outputs

The output of that sprint should be a ranked portfolio:

- **Launch now**
- **Run one more proof test**
- **Reject**

That is the practical way to validate Leon's idea. The evidence does not prove an academic law called "inevitable business," but it does support a strong product and investment heuristic: the best recovery businesses sit where failures are common and not the user's fault, remedies are enforceable, and the claim path can be turned into an almost effortless action. [filecite turn0file0](#) 28

1 3 6 7 28 Customer Evaluations of Service Complaint Experiences: Implications for Relationship Marketing - Stephen S. Tax, Stephen W. Brown, Murali Chandrashekar, 1998

<https://doi.org/10.1177/002224299806200205>

2 10 15 17 AirHelp - #1 Air Passenger Rights Experts

<https://www.airhelp.com/en/>

4 9 13 16 27 Regulation - 261/2004 - EN - Flight Compensation Regulation - EUR-Lex

<https://eur-lex.europa.eu/eli/reg/2004/261/oj>

5 8 14 Red tape

https://pt.wikipedia.org/wiki/Red_tape?utm_source=chatgpt.com

11 24 Automatic compensation: What you need to know

<https://www.ofcom.org.uk/phones-and-broadband/service-quality/automatic-compensation-need-know>

12 Airline Information for Download | Bureau of Transportation Statistics

<https://www.bts.gov/topics/airlines-and-airports/airline-information-download>

18 Claim your Right for Flight Delay Compensation - Flightright

<https://www.flightright.com/>

19 U.S.C. Title 15 - COMMERCE AND TRADE

<https://www.govinfo.gov/content/pkg/USCODE-2023-title15/html/USCODE-2023-title15-chap50.htm>

20 Amazon Compute Service Level Agreement

<https://aws.amazon.com/compute/sla/>

21 Consumer

<https://content.naic.org/consumer.htm>

22 Creating and editing GPTs | OpenAI Help Center

<https://help.openai.com/en/articles/8554397-creating-a-gpt>

23 Prompt engineering | OpenAI API

<https://platform.openai.com/docs/guides/prompt-engineering>

25 European airlines sit on €3.2bn in unpaid delay compensation

https://www.ft.com/content/f4e31fa8-b1fc-46ca-8d92-6d4c45e49313?utm_source=chatgpt.com

26 Consumer Complaint Database | Consumer Financial Protection Bureau

<https://www.consumerfinance.gov/data-research/consumer-complaints/>